



PRESENTED BY
Human Oversight
Systems Lab

Confirming Process, Missing Function

The Uninstrumented Oversight Layer in AI Transparency and Control Research

Governance records can certify that human oversight occurred without any indication whether human review actually happened.


The states that matter most for safety can be observationally identical in the audit trail.

The Problem

1.1 / OBSERVATIONAL IDENTITY

The contrast lives between the icon and the checkmark.

Two reviewers move through identical approval workflows. One reads the model output and weighs it. One does not. The contrast that matters for safety is not visible in either log entry.



1.2 / FLUENCY SUPPRESSION

Earlier automation produced legible error signatures. Language model outputs do not.

FLUENT MODEL OUTPUT
ILLUSTRATIVE / FABRICATED

Parasuraman and Wickens (2003) demonstrated that supervisory control degradation follows a logarithmic decay curve under sustained automation pressure, with calibration loss accelerating after the 40-hour threshold. The mechanism has been replicated across cockpit, clinical, and industrial settings, suggesting a stable cross-domain pattern in how trained operators disengage from automated systems they nominally supervise.

Illustrative example. Fabricated content for demonstration.

The prose reads accurately.

Five fabrications. No system flagged them. No reader noticed.

NO ALERT RAISED

FLUENT MODEL OUTPUT
VERIFIED

Bainbridge (1983) argued that automation of routine cognitive tasks erodes the attentional and skill-based conditions on which competent oversight depends, an observation Parasuraman and Manzey (2010) extended through their integration of automation bias and complacency findings. The phenomenon recurs across operational settings, though the precise conditions under which it manifests vary by domain.

The prose reads accurately.

Every claim is supported. The verification was effortful.


VERIFIED

This is not a claim that reviewers are negligent. It is a claim that the oversight function has been treated as a stable substrate rather than as a component requiring care and feeding.

The Degradation Pattern

2.1 / TRAJECTORY

Behavior degrades. The record stays uniform.



2.2 / COMPOUNDING STRUCTURE

Each transparency function compounds the next.

- Detection *pre-endorses*.
- Interpretability *re-explains*.
- Governance *certifies*.

Model capability and oversight fragility may scale together rather than in opposition.



Full paper



Follow up